

Global Investment Views Monitor Project

A Semi-Automated AI Market Intelligence Workflow for
Institutional Investment Views

*AI-assisted cross-asset research, institutional source
monitoring and automated reporting workflow*

By: Elias BANON

Academic Year: 2025–2026

Table of Contents

I.	Executive Summary	p.2
II.	Project Rationale	p.3
III.	Problem Statement	p.4
IV.	Workflow Architecture	p.4
V.	Source Map and Data Structure	p.5
VI.	Analytical Framework	p.6
VII.	Scheduled AI Task Logic	p.8
VIII.	Quality Controls and Reliability Framework	p.9
IX.	Limitations and Future Extensions	p.10
X.	Contribution	p.11
XI.	Conclusion	p.12
XII.	Appendix	p.13

1. Executive Summary

The Global Investment Views Monitor is a semi-automated generative AI workflow designed to transform fragmented public market information into a structured weekly investment intelligence brief. The project uses a large language model as an analytical assistant, combined with a curated source map, a structured Google Sheets database, a professional Google Docs template and a Make automation layer for document conversion, archiving and distribution. The purpose of the project is not to produce investment advice, trading signals or portfolio recommendations. Its objective is to replicate selected activities performed within investment strategy, wealth management and asset management teams. These activities include monitoring public macroeconomic and market news, extracting institutional investment views, comparing cross-asset signals and identifying changes in market narratives over time.

The workflow begins with a curated source map. This source map defines the institutional universe and acts as the project's single source of truth. Only sources and institutions explicitly listed in the database can be used as institutional evidence. This rule prevents the model from introducing unsupported institutions or presenting general market commentary as an official investment view. Each weekly run follows a controlled process. The scheduled AI task reviews the selected sources, checks publicly accessible publications, updates the Google Sheets database and generates a Google Docs draft using a fixed master template. The report covers the macroeconomic and market backdrop, cross-asset interpretation, institutional consensus and divergences, changes in institutional views, narrative shifts, key risks and reliability limitations. The workflow also includes a human validation step. The AI-generated report remains in a draft folder until its content, sources and institutional attributions have been reviewed. Once validated, the report is moved to a ready-to-send folder. Make then converts the Google Doc into a PDF, archives it in Google Drive and distributes it by email. This separation between analytical generation and operational distribution creates a clear control point before delivery. The final output is a professional weekly market intelligence brief designed to be concise, source-grounded and investment-focused. Rather than producing a series of isolated summaries, it translates public institutional content into a structured view of what major investment houses are communicating, where they agree, where they diverge and how their positions are changing over time.

Overall, the project demonstrates how generative AI can be integrated into a financial research workflow in a controlled and responsible manner. It shows that the value of a large language model does not lie in replacing investment judgement, but in accelerating information structuring, improving analytical consistency and supporting systematic market monitoring while preserving human oversight.

2. Project Rationale

The Global Investment Views Monitor was developed from a practical observation. Investment professionals do not suffer from a lack of information. They face the opposite problem. Market information is abundant, fragmented and unevenly structured. Banks, private banks, asset managers and CIO teams publish regular outlooks, weekly commentaries, allocation views, market updates and thematic notes. These publications often contain valuable insights. Yet they are dispersed across different platforms, written in different formats and updated at different frequencies. For an analyst, the challenge is not simply to read more. The challenge is to organise what matters, compare it across institutions and understand how views evolve over time.

This project was built in that context. It explores how generative AI can support the first layer of investment research. This layer is not about making investment decisions. It is about monitoring sources, extracting relevant signals, structuring information and preparing a reliable basis for market discussion. A large language model is useful in this process because it can read and structure large volumes of text quickly. However, in finance, speed is not enough. A generic AI summary can become vague, unsupported or misleading if it is not constrained by clear rules. Investment content requires source discipline, institutional attribution, traceability and human accountability. For this reason, the project was designed as a controlled AI-assisted workflow rather than a fully autonomous tool. The LLM operates within a fixed source map. The Google Sheets database stores the structured evidence. The analytical framework standardises the classification of views. The Google Docs template imposes consistency in the final report. Human validation remains required before distribution. Make is then used only after approval, for PDF conversion, archiving and email delivery.

The project is also consistent with my academic and professional trajectory. My previous work on ETFs examined how asset management is being transformed by product innovation, allocation tools, distribution models and investment infrastructure. This project extends that reflection to the information layer of asset management. It focuses on how investment views can be collected, structured and communicated more efficiently through data organisation, automation and generative AI. This continuity is important. Asset management is not only being transformed by new products. It is also being transformed by the way information is processed, compared and delivered. Investment teams must deal with a growing volume of publications, faster market reactions and increasing demand for clear and timely communication. A structured market intelligence workflow directly responds to this environment.

The rationale of the project is therefore both professional and methodological. Professionally, it replicates part of the monitoring work performed in investment strategy, wealth management and asset management teams. Methodologically, it tests whether generative AI can support recurring financial intelligence without weakening source traceability, institutional attribution or human oversight. The result is not an autonomous investment engine. It is a controlled research support system. Its purpose is to make institutional market views more comparable, traceable and usable over time.

3. Problem Statement

The central problem addressed by this project is not information access. Public investment content is widely available. The real issue is information comparability. Institutional investment views are rarely published in a standardised format. A CIO note, a private banking outlook, an asset manager commentary and a weekly market update may all contain relevant signals. Yet they often use different asset class definitions, regional breakdowns, investment horizons, levels of conviction and terminology. This creates a methodological problem. Without a common framework, it becomes difficult to compare views across institutions. It also becomes difficult to distinguish a durable institutional position from a short-term market comment. Similar language may describe different exposures. Different languages may describe similar views. Manual monitoring can identify individual publications, but it struggles to preserve consistency over time. A simple news summary does not solve the issue either. It may describe market events, but it does not clearly separate macro news, market reaction and official institutional positioning. A generic LLM summary creates another risk. It may produce fluent text without sufficient evidence, attribution or source control.

The problem is therefore to convert heterogeneous institutional publications into comparable, traceable and repeatable investment intelligence. This requires more than summarisation. It requires a fixed source universe, standardised classifications, explicit attribution rules, structured storage and human validation. The Global Investment Views Monitor was designed to address this methodological gap. It separates market context from institutional views. It standardises cross-asset interpretation. It tracks consensus, divergences and narrative shifts across a defined source universe. It also keeps the final output auditable by linking the weekly report to a structured database and a controlled validation process. In this sense, the project does not try to predict markets. It tries to improve the reliability of the information layer that supports market analysis. Its value lies in transforming dispersed institutional content into a structured and repeatable intelligence process.

4. Workflow Architecture

The Global Investment Views Monitor is built around a semi-automated architecture that separates analysis, validation and distribution. This structure is central to the project because the workflow does not rely on a single AI prompt. It combines several tools, each with a clearly defined function. The workflow follows a sequential process:

- Source map definition
- Scheduled AI-assisted source review
- Google Sheets database update
- Google Docs draft report generation
- Human validation
- Transfer to the ready-to-send folder
- Make-based PDF conversion, archiving and email distribution
- Transfer of the original document to the sent folder

The first layer is the source map. It defines the universe of sources that the system is authorised to review and use as evidence, subject to public availability and technical accessibility. This includes macro and market news sources, institutional publications and selected asset-level sources. The source map prevents the workflow from relying on an undefined or unstable information universe. It also ensures that institutional attribution remains controlled. The second layer is the scheduled AI task. The task acts as a recurring market intelligence analyst. It reviews the source map, checks selected

public sources, extracts relevant investment signals and updates the structured Google Sheets database. The large language model is not used as an autonomous decision-maker. It is used to process and structure information according to a defined analytical framework. The third layer is the Google Sheets database. This database acts as the operational backbone of the workflow. It stores source checks, weekly run logs, cross-asset views, consensus and divergences, risk indicators, narrative shifts and final report outputs. This makes the workflow more transparent and auditable. The weekly report is not detached from the underlying source review. It is linked to structured fields that can be checked and updated over time. The fourth layer is the Google Docs report template. Each weekly report is generated from a fixed master template. The template provides a professional structure for the final brief. It also imposes consistency across weekly outputs. This is important because the value of the workflow depends not only on the quality of the analysis, but also on the ability to present information in a clear and repeatable format. The fifth layer is human validation. The AI-generated report is first created as a draft. It is not sent automatically. The user reviews the document, checks the quality of the analysis and confirms that the report is ready. Only then is the file moved into the ready-to-send folder. This step preserves accountability and reduces the risk of distributing an incomplete or inaccurate output. The final layer is the Make automation scenario. Make does not perform the investment analysis. Its role is operational. Once a validated Google Doc enters the ready-to-send folder, Make detects the file, converts it into a PDF, archives the PDF in Google Drive, sends the report by email and moves the original document into the sent folder.

This architecture contains two main control boundaries. The first is a source boundary. The source map restricts the evidence available to the analytical process. The second is a distribution boundary. No document can enter the automated distribution process before human approval. These two boundaries make the workflow more reliable than a generic AI summary or a fully automated publication process. Overall, the architecture reflects the main principle of the project: generative AI is most useful in finance when it is embedded in a structured process. The LLM accelerates source review and synthesis, but the workflow remains governed by source discipline, database structure, human validation and controlled distribution.

5. Source Map and Data Structure

The source map is the foundation of the Global Investment Views Monitor. It defines the sources that the system is authorised to review and use as evidence, subject to public availability and technical accessibility. This is a key design choice. Without a fixed source universe, an AI-generated report could become inconsistent, unsupported or difficult to audit. The source map is organised around three layers of information:

- **Layer 1 - Macro and market media:** This layer covers macro and market news sources. These sources help describe the market backdrop, index movements, macroeconomic developments and major risk events. They provide context, but they are not treated as official institutional investment views.
- **Layer 2 - Official institutional investment publications:** This layer includes approved publications from private banks, asset managers, wealth management platforms and investment strategy departments. These sources are used to identify official cross-asset views, positioning signals, institutional consensus, divergences and changes in market narratives. They form the core evidence base for institutional views in the report.

- **Layer 3 - CIO and strategist qualitative signals:** This layer includes selected public commentary from CIOs, strategists, economists and investment specialists when available and relevant. These signals are more qualitative and indirect than formal institutional publications. They can help enrich the interpretation of market narratives, tone and risk perception, but they do not automatically represent formal house views unless they are explicitly published as official institutional positions.

Furthermore, the Google Sheet acts as the database of the workflow. It is not only a list of links. It is the structured environment in which the weekly analysis is organised, stored and controlled. The tabs most relevant to the analytical workflow include:

- **01_Source_Map:** defines the approved source universe and classifies each source by type.
- **03_Weekly_Run_Log:** records each weekly execution of the workflow.
- **04_Cross_Asset_Views:** structures views by asset class, region, conviction level, horizon and supporting institutions.
- **05_Consensus_Divergences:** identifies areas of agreement and disagreement between institutions.
- **06_Risk_Dashboard:** tracks the main risks, their probability, impact, trend and market relevance.
- **07_Narrative_Shifts:** records changes in tone, positioning or market interpretation compared with previous reports.
- **10_Source_Checks:** records whether sources were reviewed, accessible, outdated, missing or unavailable.
- **11_Brief_Output:** consolidates the structured elements used by the scheduled AI task to generate the Google Docs brief.
- **12_Task_Control:** manages workflow status, draft status, validation status and readiness for Make.

The main purpose of this structure is standardisation. Institutional publications are written in different styles and formats. The database converts them into a common analytical language. It standardises asset classes, horizons, conviction levels, supporting institutions and risk categories. This makes it possible to compare views that would otherwise remain dispersed. The source map and database therefore play a central role in the project's contribution. They transform the LLM from a generic summarisation tool into part of a controlled market intelligence workflow. The system does not simply generate text. It organises information, preserves attribution and creates a repeatable structure for comparing institutional investment views over time.

6. Analytical Framework

The analytical framework is designed to transform heterogeneous market publications into a common investment intelligence format. Its purpose is not only to summarise information but also to make institutional views comparable across sources, asset classes and time. The framework is built around a clear distinction between three types of information: market context, institutional views and analytical interpretation. Market context describes what happened during the week. Institutional views capture what approved investment houses publicly communicate. Analytical interpretation connects these inputs to cross-asset implications, consensus formation, divergences, risks and narrative shifts. This distinction is essential. A market news source can explain why equity indices moved, why bond yields

changed or why commodities reacted to a macro event. However, it cannot be treated as an official institutional allocation view. The framework therefore uses several analytical layers:

- **Macro and market backdrop:** identifies the main macroeconomic, policy, geopolitical and market developments of the week.
- **Market reaction:** records relevant moves across major equity indices, rates, credit, currencies and commodities when data can be independently validated.
- **Institutional views:** extracts public investment views from approved institutional sources only.
- **Cross-asset interpretation:** standardises explicit institutional signals across asset classes, regions, sectors, styles, duration, credit and currencies. Any implication not directly stated by the source must be identified as analytical interpretation rather than institutional evidence.
- **Consensus and divergences:** compares whether institutions are aligned, partially aligned or expressing different views.
- **Narrative shifts:** tracks changes in tone, conviction or positioning compared with previous reports.
- **Risk assessment:** identifies the main risks, their potential market impact and the assets most exposed.
- **Reliability assessment:** flags missing sources, outdated publications, weak evidence or limited market data availability.

Each view is standardised through a set of common fields. These include asset class, region, view direction, conviction level, horizon, supporting institutions, source status and analytical comment. This makes it possible to compare publications that are originally written in different styles and formats. The framework also standardises the language used in the report. For example, views can be classified as constructive, neutral, cautious, defensive or mixed. Conviction can be classified as low, medium or high. Horizons can be described as tactical, three-month, six-month or twelve-month when the source supports it. When no horizon is explicitly stated, the field remains unspecified rather than being inferred. The same logic applies to consensus and conviction. These classifications must not be based only on wording similarity. They follow predefined decision rules presented in the reliability framework. This reduces subjectivity and makes the weekly comparison more consistent.

A key feature of the framework is that it separates evidence from interpretation. Evidence comes from approved institutional sources and market data that can be independently validated. Interpretation comes from the structured comparison of that evidence. This separation reduces the risk of presenting the model's synthesis as an unsupported market view. The framework also supports longitudinal monitoring. Because views are stored in a structured format, the workflow can track whether institutions are becoming more constructive, more defensive or more divided over time. This is important because investment intelligence is not only about knowing what institutions say in a given week. It is also about understanding how their views change.

Overall, the analytical framework is the methodological core of the project. It converts fragmented publications into a standardised intelligence structure. It allows the LLM to support market monitoring while preserving comparability, attribution and analytical discipline.

7. Scheduled AI Task Logic

The scheduled AI task is configured as the recurring analytical engine of the workflow. It is designed to perform the first stage of weekly market intelligence work: reviewing sources, extracting relevant signals, updating the structured database and preparing a draft report. The task begins with the Google Sheets database. It uses the source map and task control tabs to identify which sources should be reviewed during the weekly run. This ensures that the analysis starts from the approved source universe rather than from an open-ended search. The task is therefore constrained from the beginning by the structure of the workflow. The task then reviews publicly available and technically accessible sources. It checks whether each source is available, recent and relevant to the weekly report. If a source is inaccessible, outdated or missing, the task must flag it rather than replace it with unsupported information. This rule is important because the reliability of the output depends not only on what is included, but also on what is transparently excluded.

After the source review, the task is configured to update the Google Sheets database. It records source checks, weekly run information, cross-asset views, consensus and divergences, risk indicators, narrative shifts and the structured fields required for the final brief. The database update comes before report generation. This order is intentional. It ensures that the written report is based on structured information rather than on an isolated generative output. The task then generates a Google Docs draft report from the fixed master template. The template is duplicated for each weekly run so that the original structure and formatting are preserved. The generated report remains in draft status and is stored in the draft folder. It is not sent automatically and it is not moved into the ready-to-send folder by the AI task. The task follows strict content rules. It cannot introduce institutions that are not listed in the source map. It cannot treat macro news as an official institutional view. It cannot invent figures, dates, market moves or source references. It must distinguish between confirmed evidence, analytical interpretation and missing information. When market data cannot be validated, the task must state that the figure is not available and flag it for human verification. The task also includes a rule for partial runs. A partial run may still generate a draft, but all missing sources and unvalidated fields must remain explicitly flagged for human review. The workflow must not treat an incomplete source review as a complete market intelligence assessment. The scheduled AI task also follows a formatting logic. The final report must preserve the structure of the master template, maintain a professional tone and avoid unnecessary complexity. The writing should be concise, source-grounded and investment-focused. The objective is to produce a report that resembles an institutional market intelligence note rather than a generic AI summary.

The final output of the task is therefore a draft, not a final publication. Human review remains necessary before distribution. This design keeps the LLM in the role of an analytical assistant. It accelerates source monitoring and report preparation, but it does not control the final release of the document. A weekly run is considered complete only when the source checks and structured database fields have been updated and a dated draft report has been created. Completion does not imply approval for distribution. Approval only occurs after human validation and manual transfer of the document into the ready-to-send folder. The scheduled task is central to the project because it gives the workflow its recurring nature. It turns the process from a one-off report into a repeatable weekly system. Its value lies in combining automation with structure: the task can process information quickly, but it operates within predefined boundaries, structured data fields and human validation rules.

8. Quality Controls and Reliability Framework

The Quality Controls and Reliability Framework is designed to ensure that the Global Investment Views Monitor remains transparent, auditable and source-grounded. This is essential because the project uses a large language model in a financial research context. In this setting, fluent writing is not sufficient. The output must be supported by identifiable sources, clear attribution rules and explicit limitations.

The reliability of the workflow starts with a fixed source universe. The scheduled AI task does not operate from an open-ended search process. It starts from a predefined source map, which defines the institutions, publications and market data sources that can be reviewed. This prevents the model from introducing unapproved institutions or using an unstable evidence base. Each institutional view must be linked to an identifiable and approved source. The framework also separates market context from institutional evidence. Macro news sources can be used to describe the market backdrop, index movements, policy developments or risk events. However, they cannot be presented as official asset allocation views. Institutional views must come from approved investment houses, CIO publications, private banking outlooks or asset manager commentaries. This distinction protects the report from unsupported attribution. A further control is the separation between evidence and interpretation. Evidence refers to what can be traced back to an approved institutional source or to independently validated market data. Interpretation refers to the structured comparison of this evidence across asset classes, regions, sectors, horizons and risks. Any implication not directly stated by a source must be identified as analytical interpretation rather than institutional evidence.

Source freshness is assessed relative to each source's expected publication frequency. A macro news item is expected to be recent within the weekly review period. A weekly commentary should normally correspond to the latest available edition. A monthly outlook may remain valid during the relevant month. A quarterly outlook may remain usable if it is still the latest strategic publication available. When a source is older than expected or may no longer reflect current market conditions, this must be flagged in the report. Market data validation follows the same logic of caution. Market moves can only be included when they can be independently verified through reliable public data. If an index level, weekly variation or YTD performance cannot be validated, the field must remain unavailable and be flagged for human verification. The workflow must not invent figures to complete the report. The classification system is also governed by explicit rules. Views are classified using predefined categories such as constructive, neutral, cautious, defensive or mixed. Conviction levels can be classified as low, medium or high. Horizons can be tactical, three-month, six-month or twelve-month when explicitly supported by the source. When no horizon is stated, the field remains unspecified rather than being inferred. This prevents the workflow from creating artificial precision. The reliability of each weekly run can be assessed through practical indicators. Source coverage measures the proportion of expected core sources successfully reviewed, with 70% used as a minimum reference for standard publication. Attribution coverage measures whether institutional claims are linked to approved sources, with a target of 100%. Attribution accuracy is assessed during human review by checking whether sampled claims faithfully reflect the original sources. Evidence support measures whether material claims in the report are backed by identifiable evidence. Missing sources, unvalidated figures and classification inconsistencies are recorded and disclosed rather than ignored.

Human validation remains the final control before distribution. The scheduled AI task produces a draft, not a final publication. The report remains in the draft folder until it has been reviewed. Distribution only becomes possible after the user manually moves the validated document into the ready-to-send folder. This creates a clear boundary between analytical preparation and final approval. The framework also addresses partial runs. A partial run may still generate a draft report, but all

missing sources and unvalidated fields must remain explicitly flagged. An incomplete source review must not be presented as a complete market intelligence assessment. This reliability framework is central to the credibility of the project. It shows that the workflow is not designed to maximise automation at any cost. It is designed to combine generative AI with source discipline, structured evidence, standardised classifications and human oversight.

9. Limitations and Future Extensions

The Global Investment Views Monitor is a controlled market intelligence workflow. Nevertheless, it has several limitations. These limitations define the scope of the project and prevent the output from being interpreted as investment advice, a complete market positioning tool or a fully automated research system.

The workflow depends on publicly accessible and technically available sources. Some institutional publications may be behind paywalls, temporarily unavailable or only partially accessible. When this happens, the system must flag the limitation rather than replace the missing information with unsupported content. The project also depends on the publication rhythm of each institution. Not all sources are updated weekly. Some are monthly, quarterly or linked to specific market events. This means that a weekly report may combine fresh market commentary with older but still relevant strategic publications. The report must therefore distinguish between current market developments and institutional views that may have a longer publication cycle. Another limitation is interpretation risk. Institutional language can be nuanced. A source may be cautious on equities overall but constructive on selected regions, sectors or styles. Another source may be positive on risk assets over a medium-term horizon while defensive in the short term. The workflow reduces ambiguity through standardised classifications, but human review remains necessary to avoid oversimplification. Market data validation is also a constraint. Weekly market reactions require reliable data on index levels, performance periods, closing dates and calculation conventions. If market data cannot be verified consistently, it should not be forced into the report. The correct response is to mark the field as unavailable and flag it for review. The use of a large language model creates additional reliability constraints. A LLM can structure and summarise information efficiently. However, it can produce unsupported statements if it is not constrained. This is why the project uses a fixed source universe, attribution rules, structured database fields and human validation. The LLM is treated as an analytical assistant not as an autonomous research authority. The fixed source universe improves consistency but it also creates coverage bias. The consensus identified by the workflow reflects the selected institutions only. It should not be interpreted as a complete measure of market-wide positioning. The output is therefore a structured view of the approved source universe, not a full representation of the entire investment industry. The workflow also analyses public communication rather than internal institutional decision-making. Public outlooks, CIO notes and market commentaries do not necessarily capture the full range of an institution's internal research, portfolio positioning or client-specific recommendations. The report therefore analyses published views, not complete institutional investment processes. There is also an operational dependency on several tools. The workflow relies on Google Sheets, Google Docs, Google Drive, ChatGPT scheduled tasks and Make. Changes in connector access, file permissions, document formatting or automation settings may affect execution. The system therefore requires periodic monitoring and maintenance.

Future extensions could strengthen the workflow in several directions. The most important improvement would be dedicated market data integration, which would make index levels, asset

prices and performance calculations more reliable. A second priority would be a stronger evidence-level audit trail, including source excerpts, review dates and links for each reported view. A third priority would be formal reliability scoring based on source coverage, freshness, attribution strength and data completeness. Further extensions could include a historical dashboard to track changes in institutional views over several weeks or months, automated charts on market performance and risk trends, a broader source universe, bilingual output in English and French, and API-based integration for more robust data retrieval, document generation and monitoring.

The main limitation of the current version is also one of its main strengths. The workflow remains intentionally controlled. It favours reliability, traceability and human accountability over full automation. This makes it more suitable for a financial research context, where the quality of evidence matters as much as the speed of production.

10. Contribution

The contribution of the Global Investment Views Monitor is to transform generative AI from a generic summarisation tool into a controlled investment intelligence workflow. The project does not simply ask a large language model to summarise market publications. It embeds the model within a structured process built around a fixed source universe, a standardised analytical framework, a database, a reporting template, human validation and controlled distribution.

The main methodological contribution is the conversion of heterogeneous institutional publications into comparable, traceable and consistently structured investment intelligence over time. Public market outlooks are difficult to compare because institutions use different formats, vocabularies, horizons and levels of conviction. The workflow addresses this problem by converting these publications into common analytical fields while preserving the distinction between institutional evidence and analytical interpretation. The project also contributes by clarifying the role that generative AI can play in a financial research environment. The LLM is not presented as an investment decision-maker. It is used as a structured research assistant. Its role is to accelerate source review, organise fragmented information, support cross-asset synthesis and prepare a draft report. Final judgment remains human, and the distribution process only begins after validation. Another contribution lies in the workflow design. The project combines flexible AI-assisted analysis with deterministic operational automation. ChatGPT supports the analytical stage, while Google Sheets provides structured storage, Google Docs standardises reporting and Make manages document conversion, archiving and distribution after approval. This separation is designed to improve robustness compared with an isolated prompt-based process by introducing structured inputs, explicit controls and a validation boundary. The project also has a practical contribution for investment communication. It produces a weekly brief designed to be readable, structured and decision-useful without becoming a recommendation engine. It helps identify what selected institutions are saying, where they agree, where they diverge and which risks or narrative shifts deserve attention.

From an academic perspective, the project extends my previous work on the transformation of asset management. My ETF research focused on investment products, allocation tools and industry structure. This project focuses on the information infrastructure that supports investment analysis. Both projects examine how finance is being reshaped by standardisation, technology, automation and new forms of investment communication. As a result, this project is therefore not the creation of a forecasting model. It is the design of a controlled framework for recurring market intelligence. Its

distinctive contribution is not the use of a large language model alone, but the design of the controls, data structures and validation boundaries that make its use credible in a recurring financial research process.

11. Conclusion

The Global Investment Views Monitor demonstrates how generative AI can be used responsibly in a financial research workflow. The project starts from a clear problem. Investment information is abundant but institutional views are difficult to compare, trace and monitor over time. It responds to this problem by building a semi-automated process that combines AI-assisted analysis, structured data, professional reporting and controlled distribution. The implementation produced a repeatable workflow structure linking approved sources, structured analytical fields, a standardised report and a controlled distribution sequence. Its architecture combines a restricted evidence universe, structured analytical storage, standardised reporting and a mandatory validation boundary before distribution. This makes the system more disciplined than a generic AI summary and more transparent than an uncontrolled automated report. The project's value lies in the way it defines the role of the LLM. The objective is not to remove human judgment from the research process. It is to improve the quality, speed and consistency of the information layer that supports that judgment. A large language model can be useful when it is guided by clear rules, reliable sources and a defined analytical structure. The project also shows that automation and control are not opposite objectives. Full automation may increase speed, but it can also increase the risk of unsupported output. A controlled workflow can reduce repetitive manual work while preserving traceability, attribution and accountability. This balance is particularly important in asset management, where the credibility of a report depends on the quality of its evidence as much as on the clarity of its writing.

The current version remains a project-scale implementation. It depends on public sources, manual validation and the technical availability of connected tools. However, these limitations are also part of its design discipline. The workflow does not claim to replace professional research systems, internal investment committees or portfolio construction tools. It focuses on a specific and realistic use case: recurring, source-grounded market intelligence. Overall, the Global Investment Views Monitor provides a practical demonstration of disciplined AI augmentation in investment research. It transforms scattered public information into a structured weekly intelligence process while preserving source discipline, human review and methodological transparency. The project suggests that the most credible use of generative AI in investment research is not autonomous decision-making but disciplined augmentation. In other words, using technology to improve the processing and communication of evidence while preserving transparency and human responsibility.

12. Appendix


12.A Source Universe

The source universe defines the approved information base of the Global Investment Views Monitor. It is organised into three layers. Each layer has a specific role in the workflow and a specific treatment in the final report.

Layer 1 - Macro and market media: includes macro and market media sources used to describe the weekly market backdrop. These sources help identify major macroeconomic developments, policy events, index movements, geopolitical risks and market reactions. They provide context for the report, but they are not treated as official institutional investment views. Representative sources include Reuters Markets, Bloomberg Markets, Financial Times Markets, The Wall Street Journal Finance, CNBC Markets, The Economist Finance and Economics, MarketWatch and Yahoo Finance.

Layer 2 - Official institutional investment publications: includes approved publications from private banks, asset managers, wealth management platforms and investment strategy departments. These sources are used to identify official investment views, cross-asset positioning signals, institutional consensus, divergences and changes in market narratives. They form the core evidence base for institutional views in the report. Representative sources include UBS CIO publications, HSBC Private Banking investment views, J.P. Morgan Private Bank market insights, Goldman Sachs Private Wealth Management insights, Julius Baer market outlooks, Citi Wealth or Citi Private Bank market publications and BlackRock Investment Institute commentary.

Layer 3 - CIO and strategist qualitative signals: includes selected public commentary from CIOs, strategists, economists and investment specialists when available and relevant. This layer is more qualitative and indirect than formal institutional publications. It can help enrich the interpretation of market narratives, tone, risk perception and changes in emphasis. However, these signals do not automatically represent formal house views unless they are explicitly published as official institutional positions.



	A	B	C	D	E	F	G	H
	Scope	Layer	Institution / Source	Source name	Source type	Frequency	Reliability tier	Access status
1	Core	Layer 1	Reuters	Reuters Markets	News / markets	Weekly	Tier 2	Public / partial
2	Core	Layer 1	Bloomberg	Bloomberg Markets	News / markets / market data	Weekly	Tier 2	Subscription / partial
3	Core	Layer 1	Financial Times	FT Markets	News / analysis	Weekly	Tier 2	Subscription / partial
4	Core	Layer 1	Wall Street Journal	WSJ Finance / Markets	News / analysis	Weekly	Tier 2	Subscription / partial
5	Core	Layer 1	CNBC	CNBC Markets	News / markets	Weekly	Tier 2	Public / partial
6	Core	Layer 1	The Economist	Finance and economics	News / analysis	Weekly / event-driven	Tier 2	Subscription / partial
7	Core	Layer 1	MarketWatch	MarketWatch Markets	News / markets	Weekly / event-driven	Tier 3	Public / partial
8	Core	Layer 1	Yahoo Finance	Yahoo Finance Markets / News	News / market data	Weekly / event-driven	Tier 3	Public / partial
9	Core	Layer 2	UBS Global Wealth Management	UBS CIO House View	Official CIO view / house view	Daily / weekly / monthly	Tier 1	Public / partial
10	Core	Layer 2	UBS Global Wealth Management	UBS Weekly Key Messages	Official weekly CIO commentary	Weekly	Tier 1	Public / partial
11	Core	Layer 2	HSBC Private Bank	HSBC Investment Outlook Programme	Official investment outlook	Quarterly / event-driven	Tier 1	Public / partial
12	Core	Layer 2	HSBC Private Bank	HSBC House Views	Official house views	Regular / event-driven	Tier 1	Public / partial
13	Core	Layer 2	J.P. Morgan Private Bank	J.P. Morgan Private Bank Markets & Investing	Official private bank insights hub	Weekly / event-driven	Tier 1	Public / partial
14	Core	Layer 2	J.P. Morgan Private Bank	Global Investment Strategy View	Official strategy view	Weekly / monthly / event-driven	Tier 1	Public / partial
15	Core	Layer 2	J.P. Morgan Private Bank	Top Market Takeaways	Official weekly update	Weekly	Tier 1	Public / partial
16	Core	Layer 2	Goldman Sachs Private Wealth Management	Investment Strategy Group	Official PWM strategy hub	Event-driven	Tier 1	Public / partial
17	Core	Layer 2	Goldman Sachs Private Wealth Management	Economic and Financial Market Outlook	Official PWM outlook	Annual / event-driven	Tier 1	Public / partial
18	Core	Layer 2	Goldman Sachs Private Wealth Management / Ayco	Quarterly Market Update	Official market update	Quarterly	Tier 1	Public / partial
19	Core	Layer 2	Julius Baer	Julius Baer Market Outlook	Official market outlook	Annual / mid-year / event-driven	Tier 1	Public / partial
20	Core	Layer 2	Julius Baer	Julius Baer Market Insights	Official market insights	Regular / event-driven	Tier 1	Public / partial
21	Core	Layer 2	Julius Baer	Julius Baer Mid-Year Market Outlook	Official market outlook / report	Mid-year	Tier 1	Public / partial
22	Core	Layer 2	Citi Wealth / Citi Private Bank	Citi Global Market Insights	Official wealth management insights	Weekly / event-driven	Tier 1	Public / partial
23	Core	Layer 2	Citi Wealth	Citi Weekly Market Update	Official weekly market commentary	Weekly	Tier 1	Public / partial
24	Core	Layer 2	Citi Wealth	The Short and Long — Macro Investment View	Official macro investment view / PDF	Quarterly	Tier 1	Public / partial
25	Core	Layer 2	BlackRock Investment Institute	BlackRock Weekly Market Commentary	Official investment commentary	Weekly	Tier 1	Public / partial

12.B Google Sheets Database Structure

The Google Sheets file acts as the operational database of the workflow. It is not only a list of links. It stores the source map, weekly source checks, structured analytical fields, risk indicators, narrative changes and final brief inputs. The main tabs are:

- **01_Source_Map**: defines the approved source universe and classifies each source by layer, institution, type and role in the workflow.
- **02_Analytical_Framework**: records the standard categories used to classify views, including asset classes, regions, conviction levels, horizons, risk types and source status.
- **03_Weekly_Run_Log**: tracks each weekly execution of the workflow, including run date, report period, completion status, missing sources and validation status.
- **04_Cross_Asset_Views**: stores standardised views by asset class, region, direction, conviction, horizon, supporting institutions and analytical comments.
- **05_Consensus_Divergences**: identifies areas where institutions are aligned, partially aligned or expressing different views.
- **06_Risk_Dashboard**: tracks the main risks identified during the weekly review, including probability, impact, trend, exposed assets and monitoring triggers.
- **07_Narrative_Shifts**: records changes in tone, conviction or positioning compared with previous reports.
- **08_Prompt_Library**: stores the main task instructions, quality rules, reporting instructions and control prompts used by the workflow.
- **09_CIO_Watchlist**: tracks selected public CIO or strategist signals when relevant. These signals are treated carefully and do not automatically represent formal institutional house views.
- **10_Source_Checks**: records whether each source was reviewed, accessible, recent, outdated, missing or requiring human verification.
- **11_Brief_Output**: consolidates the structured elements used by the scheduled AI task to generate the Google Docs weekly brief.
- **12_Task_Control**: manages workflow status, draft status, validation status and readiness for Make distribution.

12.C Analytical Classification Rules

The analytical classification rules are designed to make institutional views comparable across sources and over time. They reduce ambiguity and prevent the workflow from relying only on wording similarity.

View direction

- **Constructive**: the source expresses a positive or supportive view on an asset class, region, sector or theme.
- **Neutral**: the source does not express a strong positive or negative preference.
- **Cautious**: the source expresses concern, selectivity or reduced confidence.
- **Defensive**: the source favours risk reduction, protection or safer assets.
- **Mixed**: the source expresses different views across horizons, regions or segments.

Conviction level

- **High**: the source expresses a clear and repeated preference, supported by explicit reasoning.
- **Medium**: the source expresses a visible view, but with some qualification or uncertainty.
- **Low**: the source expresses a limited, tentative or weakly developed view.

- **Not stated:** the source does not provide enough information to classify conviction.

Investment horizon

- **Tactical:** short-term positioning or near-term market view.
- **3M:** view expressed over an approximate three-month horizon.
- **6M:** view expressed over an approximate six-month horizon.
- **12M:** view expressed over a one-year horizon.
- **Unspecified:** no explicit horizon is provided by the source.

Source status

- **Reviewed:** the source was checked during the weekly run.
- **Accessible:** the source could be accessed and reviewed.
- **Inaccessible:** the source could not be accessed during the run.
- **Outdated:** the source was accessible but older than expected given its normal publication frequency.
- **Missing:** the expected publication could not be found.
- **Requires human verification:** the source or data point requires manual review before being used confidently.

12.D Reliability Indicators

The reliability indicators are used to assess the quality of each weekly run. They do not guarantee that the report is complete. Nevertheless, they make its coverage, evidence base and limitations more visible:

- **Source Coverage** measures the proportion of expected core sources that were successfully reviewed. A minimum reference threshold of 70% is used for standard publication.
- **Attribution Coverage** measures the proportion of institutional claims linked to approved sources. The target is 100%, because institutional views should not appear in the report without source linkage.
- **Attribution Accuracy** is assessed during human review. It checks whether sampled claims faithfully reflect the content of the original source.
- **Evidence Support** measures whether material claims in the report are backed by identifiable evidence. The target is 100% for material claims used in the final brief.
- **Data Freshness** measures whether reviewed sources fall within their expected publication window. This is assessed relative to each source's normal publication frequency.
- **Missing Sources** records the number of inaccessible, outdated or unavailable sources. These gaps must be disclosed rather than ignored.
- **Unvalidated Data** records the number of market figures or performance data points requiring human verification.
- **Classification Consistency** measures whether views are classified using the predefined categories of the analytical framework.
- **Time Efficiency** compares the estimated manual monitoring time with the AI-assisted workflow time after comparable runs.

12.E Scheduled AI Task Rules

The scheduled AI task acts as the recurring analytical engine of the workflow. Its role is to review sources, update the structured database and prepare a draft report.

Prompt Summary: *The task follows several core rules. The task must start from the approved source map. It must not introduce institutions that are not listed in the database. The task must update the Google Sheets database before generating the Google Docs draft. The report must be based on structured fields rather than on an isolated generative output. The task must distinguish between market context, institutional views and analytical interpretation. Macro news can explain market developments, but it cannot be presented as an official institutional view. The task must not invent figures, dates, source references, market moves or institutional positions. If a figure cannot be validated, it must remain unavailable and be flagged for human verification. The task must flag missing, outdated or inaccessible sources. It must not silently replace them with unsupported information. The task must preserve the Google Docs master template. Each weekly report should be generated from a duplicated version of the template, not from a blank document. The task must create a draft only. It must not move the report into the ready-to-send folder. It must not send the report. It must not trigger the distribution process. A weekly run is complete only when the source checks and structured database fields have been updated and a dated draft report has been created.*

Tous les lundis

Weekly Views Monitor

PROMPT

TASK TITLE
Weekly Investment Views Database and Report Generator

SCHEDULE
Every Monday at 08:00 Europe/Paris

MISSION
Execute the Weekly Investment Views Database and Report Generator using the

FRÉQUENCE

Répéter	Personnalisé ▾
Fréquence	Hebdomadaire ▾
Jours de la semaine	lun. ▾
Heure	8:00 ▾

12.F Make Automation Scenario

Make is used as the operational automation layer of the project. It does not perform investment analysis. It is only activated after human validation. The Make scenario follows this sequence:


- Google Drive watches the folder **03_Ready_To_Send**.
- A filter checks that the file name contains **Global_Investment_Views_Monitor**.
- Google Drive downloads the validated Google Doc as a PDF.
- Google Drive uploads the PDF to **04_PDF_Archive**.
- Gmail sends the PDF as an email attachment.
- Google Drive moves the original Google Doc to **06_Sent_Docs**.

This process creates a clear distribution boundary. The AI task can generate a draft report, but it cannot send it. Distribution only begins when the user manually moves a reviewed document into the ready-to-send folder. The Make automation therefore supports operational consistency without weakening human control. It ensures that validated reports are converted, archived and distributed in a repeatable way.

Mon Drive > GLOBAL_INVESTMENT_... ▾

Type ▾ Contacts ▾ Date de modification ▾ Source ▾

Nom ▾	Propriétaire	Date de modification	Taille du fich	Tri
05_Sent_Docs		14:10 moi	—	⋮
04_PDF_Archive		28 juil. moi	—	⋮
03_Ready_To_Send		28 juil. moi	—	⋮
02_Drafts		28 juil. moi	—	⋮
01_Source_Map		28 juil. moi	—	⋮
00_Template		28 juil. moi	—	⋮



12.G Example Output Snapshot

The example weekly brief illustrates the final output of the workflow. It shows how the structured database and analytical framework are converted into a professional market intelligence report. The example output includes the following sections:

- Executive View

- The purpose of this appendix is not to reproduce the full weekly brief. The full report can be provided as a separate deliverable. The appendix only presents selected screenshots or extracts to demonstrate the structure, formatting and analytical logic of the final output. The example output also shows the practical result of the workflow. The project does not only define a methodology. It produces a professional weekly brief that can be archived, reviewed and distributed through a controlled process.

12.H Project Repository Structure

- **README.md**: short project overview presenting the objective, rationale, workflow architecture, tools used, quality controls and main outputs.
- **project_report.pdf**: professional report designed to present the project in academic applications, interviews or portfolio reviews.
- **example_weekly_brief.pdf**: sample output generated by the workflow, showing the final format of the weekly market intelligence brief.
- **source_map_export.xlsx**: exported version of the Google Sheets workflow database, including the source map, analytical framework, source checks, cross-asset views, consensus and divergences, risk dashboard, narrative shifts and brief output structure.

- **prompts/**: folder containing the cleaned scheduled AI task instructions, reporting rules, source discipline rules and quality control prompts used to guide the LLM.
- **templates/**: folder containing the report template structure and documentation explaining how the Google Docs master template is used for weekly report generation.